

1. (20 pts.) Problem 1

Run DFS on the tree starting from r and store the previsit and postvisit times for each node. Since the given graph is a tree, and we started at the root, the DFS tree is the same as the given tree. We would like to answer queries of the form “is u an ancestor of v ?” (or equivalently “is v a descendant of u ?”) in $O(1)$ time. For this, note that u is an ancestor of v in the DFS tree if and only if $pre(u) < pre(v) < post(v) < post(u)$. After computing all the pre and post numbers, this condition can be checked for any pair of vertices u and v in $O(1)$ time. Running DFS on a tree takes $O(|V|)$ since $|E| = O(|V|)$.

2. (20 pts.) Problem 2

- (a) This setting can be represented as a directed graph. We view the intersections as its vertices with the streets being directed edges, since they are one-way. Then the claim is equivalent to saying that all the West vertices are in a SCC, and all the East vertices are in a SCC (potentially the same one). This is true if and only if the graph has only two or one SCCs. Say there are n nodes in West, and m nodes in East. Since all of East or West needs to be in the same SCC, if there are at least 2 SCCs, each SCC will have either n or m nodes, so there cannot be a third SCC. First, run the SCC algorithm, if there is only 1 SCC, the answer is YES. If there are two SCCs, pick one and isolate it from the other, then run Explore from any node in it. If all the nodes in the SCC are West, or all the nodes are East, the answer is YES, otherwise it is no. If there are not 1 or 2 SCCs, the answer is NO. Explore and the SCC algorithm take linear time.
- (b) The weaker claim says that starting from the park in either district, one cannot get to any other SCC in the graph. This is equivalent to saying that the SCC containing the vertices corresponding to each park are each in (potentially different) sink components. Regarding how the claim can be checked in linear time, there are two possible answers:
 - (1) First find all the SCCs by running the SCC algorithm, and then run Explore from the vertex corresponding to the West park (in the original graph). If any vertices visited through this Explore procedure are not in the SCC that the park belongs to, then this SCC is not a sink SCC. Otherwise, it is. The running time is linear because the running time of both the SCC algorithm and Explore is linear, as we know from class. Doing the same for the East park takes linear time as well.
 - (2) Run the SCC algorithm. Note that this algorithm for finding SCCs progressively discards sinks SCCs one by one. A component found by the algorithm is a sink if and only if there are no edges going out of the component. We use this to check if the SCCs containing each park vertex are sink components. The running time is linear because the running time of the SCC algorithm is linear.

3. (20 pts.) Problem 3

1. **First solution:** Let us call a vertex from which all other vertices are reachable, a vista vertex. If the graph has a vista vertex, then it must have only one source SCC (since two source SCCs are not reachable from each other), which must contain the vista vertex (if it's in any other SCC, there is no path from the vista vertex to the source SCC). Moreover, in this case every vertex in the source SCC will be a vista vertex. The algorithm is then simply to a DFS starting from any node and mark the vertex with the highest post value. This must be in a source SCC. We now run Explore from this

vertex to check if we can reach all nodes. Since the algorithm just uses decomposition into SCCs and Explore, the running time is linear.

2. **Second solution:** First observe that if such vertex exists, then it should lie on one of the source SCCs, since otherwise, we will have a cycle in corresponding DAG created by SCCs which is a contradiction. Second, if there are more than one source SCC, then vertices on these source SCCs can not reach each other since the in-degree of each source SCC is zero. Thus in this case, we do not have vertex from which all vertices are reachable. So, we can say such vertex exists if and only if there is only one source SCC, and in this case we return a node from the source SCC. However, note that we can run the SCC algorithm discussed in class to find the SCCs, and its corresponding DAG in linear time. Then, if there were at least two nodes on DAG with in-degree of zero, algorithm returns "such vertex does not exist", and otherwise, it returns a node from the SCC that has zero in-degree.

4. (20 pts.) Problem 4

- (a) Assume it is possible to have BFS forward edges. Assume one of this type of edges is (u, w) , which leads a node u to its non-child descendant w . As u 's descendant, w is visited after u is visited. Due to the edge (u, w) , $dist[w] = dist[u] + 1$. So, w is u 's child which contradicts with the earlier assumption that w is u 's non-child descendant. Therefore, it is impossible to have BFS forward edges.

- (b) Algorithm:

- Run BFS on G and obtain the BFS tree. The edges in the BFS tree are tree edges.
- Intuition: both BFS back edges and DFS back edges require the tail to lead to its ancestor as the head. Also, both BFS cross edges and DFS cross edges require the tail to lead to a node that's neither its ancestor nor its descendant as the head. The only difference between BFS edges and DFS edges are that one is regarding the BFS tree and the other is regarding the DFS tree. So, the key to classify BFS back edges and BFS cross edges is to determine the ancestor/descendant relationship between the tail and the head of the edges. We can make use of pre-visit number and post-visit number of each node in a tree to achieve this, and those numbers can be obtained by running DFS on the tree.
- Run DFS on BFS tree to obtain the pre-visit number and post-visit number for each vertex on the BFS tree.
- To classify the edges that are not in the BFS tree (non tree edges), for an edge (w, v) , if $pre[v] < pre[w] < post[w] < post[v]$, then it's a back edge. If $pre[v] < post[v] < pre[w] < post[w]$ or $pre[w] < post[w] < pre[v] < post[v]$, then it's a cross edge.

Run time analysis:

- BFS takes $O(|V| + |E|)$ time.
- DFS takes $O(|V| + |E|)$ time.
- Going through the edge set takes $O(|E|)$ time.

Thus, overall, the algorithm takes linear time, $O(|V| + |E|)$.

5. (20 pts.) Problem 5

- (a) Notice that in order to have a directed edge from vertex u to vertex v , the clause $\hat{u} \vee v$ needs to be in ϕ . Hence, the edge from u to v is equivalent to $u \Rightarrow v$. Consequently, a path from x to y means that $x \Rightarrow y$. If x and \hat{x} are in the same strongly connected component, then we have $x \Rightarrow \hat{x}$ and $\hat{x} \Rightarrow x$. Then there is no way to assign a value to x to satisfy both these implications.

- (b) **Consider the following algorithm:** We recursively find sinks in the graph of strongly connected components and assign true to all the literals in the sink component (this means that if a sink component contains the literal \hat{x} , then we assign $\hat{x} = \text{true}$ and $x = \text{false}$). We will then remove all the variables which have been assigned a value from the graph.

Correctness: The algorithm will output an assignment. Assume on the contrary that the assignment does not satisfy all the clauses, say, a clause $\alpha \vee \beta$ is not satisfied by the assignment, where α and β are literals. So the algorithm sets $\alpha = \text{false}$, $\beta = \text{false}$, $\hat{\alpha} = \text{true}$, $\hat{\beta} = \text{true}$. Note that the algorithm sets α and $\hat{\alpha}$ at the same time, and similarly for β and $\hat{\beta}$.

- Case 1 - the algorithm sets α before β : at the time the algorithm sets $\hat{\alpha} = \text{true}$, $\hat{\alpha}$ is in a sink component. Since β and $\hat{\beta}$ are set at a later time, the vertices for β and $\hat{\beta}$ must still exist. So, the edge $\hat{\alpha} \Rightarrow \beta$ is still in the graph just before $\hat{\alpha}$ is removed, and hence β must also be in the same sink component. So the algorithm should set $\beta = \text{true}$ at the same time it sets $\hat{\alpha} = \text{true}$, a contradiction.
 - Case 2 - the algorithm sets α and β at the same time: the algorithm sets $\hat{\alpha} = \text{true}$ and $\hat{\beta} = \text{true}$ at the same time, so they are in the same sink component. Right before they are set, the edge $\hat{\alpha} \Rightarrow \beta$ is still in the graph, so β is in the same component as $\hat{\alpha}$. So β is in the same sink component as $\hat{\beta}$, implying that they are in the same strongly connected component (at the time they are set, and hence in the original graph), a contradiction.
 - Case 3 - the algorithm sets α after β : symmetric to Case 1, a contradiction.
- (c) To perform the operations in the previous part, we only need to construct this graph from the formula, find its strongly connected components, and identify the sinks - all of which can be done in linear time.

6. (0 pts.) Acknowledgments

- (a) I did not work in a group.
(b) I did not consult with anyone other than my group members.
(c) I did not consult any non-class materials.

Rubric:

Problem 1, 20 pts

- 10 pts for a DFS from root r
- 10 pts for u is an ancestor of v if and only if $pre(u) < pre(v) < post(v) < post(u)$

Problem 2, 20 pts

- (a) 10 points
 - 4 points: correctly formulate the setting as a directed graph.
 - 3 points: correctly state what the mayor claims in graph theory terminology.
 - 3 points: provide an algorithm that can check the claim and explain why the algorithm takes linear time.
- (b) 10 points
 - 5 points: correctly state what the weaker claim says in graph theory terminology.
 - 5 points: provide an algorithm that can check this weaker claim and explain why the algorithm takes linear time.

Problem 3, 20 pts

1. First solution (20 pts):
 - Designing the correct algorithm (10 pts): 5 pts for running the first DFS, and 5 pts for running second DFS on the vertex with highest post value number which defined by first DFS.
 - Proof of correctness (10 pts): 5 pts for showing that there should be at most one source SCC, and 5 pts for saying that the highest post value number will lie on a source SCC
2. Second solution (20 pts):
 - Designing the correct algorithm (10 pts): 5 pts for running linear time algorithm for finding SCCs, and 5 pts for finding vertices with zero in-degree
 - Proof of correctness (10 pts): 5 pts for showing that there should be at most one source SCC, and 5 pts for saying that vertices with zero in-degree of corresponding DAG can give us such sources.

Problem 4, 20 pts

- (a) 6 points: provide an explanation, and it makes sense.
- (b) 14 points
 - 10 points: provide an algorithm that can correctly categorize the edges in linear time.
 - 4 points: explain why this algorithm takes linear time.

Problem 5, 20 pts

- (a) 8pts
 - 2pts for if there is an edge from u to v , there should a clause $\hat{u} \vee v$ in ϕ
 - 2pts for the path from x to y means $x \Rightarrow y$
 - 4pts for conclusion why x and \hat{x} can not be in the same SCC
- Or any reasonable explanation worth 8 pts

(b) 8pts - reasonable explanation of why the given algorithm is correct.

(c) 4pts

- 1pts for construct the graph
- 1pts for find the connected components
- 1pts for identify the sinks
- 1pts make the conclusion of linear time